



The Business School  
for the World®

# Machine Learning and Optimization

## Lecture 3

# Supervised learning: basics & regression recap

Professor Georgina Hall



# Agenda for today

## **Learning objectives:**

- Revisit linear regression and regression
- Understand some basic concepts of supervised learning: overfitting, testing/training/validation sets

## **How will we get there?**

- Make use of your prior knowledge about regression
- Activity: M. Gelato's ice-creams

# Supervised learning

Recall from Lecture 1:

What is the difference  
between **supervised** and **unsupervised** learning?

Data with **labels** or not.

→ Notion of what the “right” answer is in supervised

You already know one example of supervised learning:  
**linear regression**

# Linear Regression

# Before moving on...

Please download

- ML&O Lecture 3 - Exercise\_book.ipynb
- women\_data.csv
- Gelato\_times\_sales.csv

# Linear regression (1/7)


We are going to do regression on an easy dataset: **women\_data**

1. Take a look at it. How many observations?
2. Plot weight as a function of height using `plt.scatter`.
3. Which one is the independent variable? The dependent variable?

```
women_data.head()
```

	neight	weight
0	58	115
1	59	117
2	60	120
3	61	123
4	62	126
5	63	129
6	64	132
7	65	135
8	66	139
9	67	142
10	68	146
11	69	150
12	70	154
13	71	159
14	72	164

```
plt.scatter(women_data["height"],women_data["weight"])  
plt.xlabel("height")  
plt.ylabel("weight")
```



# Linear regression (2/7)

- Up until now, we've called **feature** any column/variable in our dataset.
- In regression, there is a special variable/feature: **the dependent variable**.
- We refer to the **dependent variable** as the **label** (=“special” feature). The **independent variables** are the **features**.


**Label here is weight. Feature is height.**

**Goal: If you are given a height, can you predict a weight?**

**How to do this? One way: linear regression.**

# Linear regression (3/7)

In linear regression, we **fit a line to the data**. How do we add this line?



**Input:**  $(height_i, weight_i)$  pairs (15 of them)

**Goal:** find numbers ( $a = \text{intercept}$ ,  $b = \text{slope}$ ) such that sum of residuals squared  $(weight_1 - a - b \cdot height_1)^2 + \dots + (weight_n - a - b \cdot height_n)^2$  is smallest possible

**Output:** Once we have  $(a, b)$ , we can obtain the predicted values  $weight_{new} = a + b \cdot height_{new}$



# Linear regression (4/7)

How to code this up in Python?

We use the **Scikit package**:

```
from sklearn.linear_model import LinearRegression
```

**First part: fitting the model, i.e., fitting the line:**

```
X=women_data[["height"]]  
Y=women_data[["weight"]]  
  
lm = LinearRegression().fit(X, Y)
```

Specify what  
your x and y  
variables are

Model you're  
interested in

Fits that model to your  
data

# Linear regression (5/7)

## Second part: getting relevant parameters outputted

```
print("Intercept = ",lm.intercept_) # Print the resultant model intercept  
print("Model coefficients = ", lm.coef_) # Print the resultant model coefficients (in order of variables in X)  
print("R^2 =",lm.score(X,Y)) # Print the resultant model R-squared
```

See everything you can get in the documentation:

[https://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.LinearRegression.html](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LinearRegression.html)

## Third part: getting the predictions

```
Y_pred=lm.predict(X)
```


Model is named “lm” and we predict from that model

How are we getting these predictions? If we had a new datapoint, e.g. 62.5, how would it work to get a prediction?

# Linear regression (6/7)

## Fourth part (Optional): Plotting

How can we use `plt.plot()` and `plt.scatter()` to get the graph below?



```
plt.scatter(X,Y)
plt.plot(X,Y_pred,c="red")
plt.xlabel("height")
plt.ylabel("weight")
```


# Linear regression (7/7)

Why do we do linear regression?

The goal is to predict new values.

For example, if an American woman ages 30-39 gives us her height, we should be able to infer her weight by using our regression line.

Linear regression is not always good enough....




# Polynomial Regression

# Regression in general

- In (general) regression, our goal is to fit a **curve** to data (not just a line as in linear regression).
- Of particular interest to us now is going to be **polynomial regression**.

Why? It is “easy” to do like linear regression but is more versatile.



Source: <http://brandon.ai/>

# Polynomial Regression (1/5)

We use scikit-learn again for this, but it is a bit more complicated.

**Let's take a look at the code together!**

**First part: specify the degree of interest and the features and label.**

```
degree=3  
X=women_data[["height"]]  
Y=women_data[["weight"]]
```

# Polynomial Regression (2/5)

**Second step: go from the datapoints to their “polynomial version”**

For example, height datapoint 58 becomes  $[1, 58, 58^2, 58^3]$

Intercept      Powers up to 3  
as degree=3

```
poly = PolynomialFeatures(degree) #define the polynomial  
X_poly=poly.fit_transform(X) #map all the values of X as [1,x,x^2,x^3, etc]
```

**Your turn!**

**Double-check that this is the case by taking a look at X and X\_poly.**

$[1.00000e+00, 5.80000e+01, 3.36400e+03, 1.95112e+05]$

=1

=58

=  $58^2$

=  $58^3$



# Polynomial Regression (3/5)

Third step: fit a Linear Regression model to the polynomial datapoints.

```
polyreg = LinearRegression().fit(X_poly, Y)
```

## Why linear regression?

We want to find the **coefficients** of the polynomial:

$$c_0 + c_1x + c_2x^2 + c_3x^3$$

Now that we've made the datapoints polynomial, the expression above is **linear** in the coefficients  $\Rightarrow$  **Linear Regression**

# Polynomial Regression (4/5)

**Fourth step: get the coefficients ( $c_0, c_1, c_2, c_3$ ) and  $R^2$ .**

**Your turn!**

```
print(polyreg.coef_) #print these coefficients  
print(polyreg.score(X_poly,Y)) #print R^2
```

```
[[ 0.00000000e+00  4.64107891e+01 -7.46184371e-01  4.25255572e-03]]  
0.9997816939979363
```

**Fifth step: Predict new points**

```
y_pred=polyreg.predict(X_poly)
```

Same as before, except that we're applying it to  $X\_poly$ , the transformed variables.

# Polynomial Regression (5/5)

**Sixth step: plot!** A bit more complicated than for linear regression:


- A line is defined by two points but not a polynomial
- We need many more points to be able to generate a “pretty” curve.

```
linepoints = np.linspace(np.min(X), np.max(X), 100)
linepoints_poly=poly.fit_transform(linepoints)
linepoints_pred=polyreg.predict(linepoints_poly)

plt.plot(linepoints,linepoints_pred)
```

**Here 100  
points  
between 58  
and 72!**

**Your turn! Try running this code!**



# Discovering the basic concepts of supervised learning

Let's start with an activity in  
BORs.

# Selling ice-creams



M. Gelato is the owner of an ice-cream parlor.

- Over the course of the year, he has written down some of his **daily sales**.
- He knows that his sales are **periodic**: i.e., every year, he sells in approximately the same way
- He would like to fit a curve to his data so as to better predict what his sales are going to be next year.

**Goal: help M. Gelato!**

In BORs (LG 1-15) with your groups: complete part 4 of the Notebook.

Slides are on the course website.

# Activity wrap-up (1/3)


```
Gelato=pd.read_csv("Gelato_Times_Sales.csv")
Gelato
```

	Times	Sales
0	1	122
1	29	215
2	60	228
3	88	225
4	121	190
5	150	145
6	184	125
7	218	43
8	250	12
9	281	4
10	318	32
11	350	80
12	365	110


## Why do we normalize the data here?

We build  $[1, x, x^2, \dots, x^{12}]$  in the worst case and  $365^{12}$  is huge! If we normalize, we don't have as huge numbers.


# Activity wrap-up (2/3)




$d=1$ ,  $R^2=0.516$




$d=2$ ,  $R^2=0.531$




$d=3$ ,  $R^2=0.979$



$d=5$ ,  $R^2=0.985$



$d=8$ ,  $R^2=0.992$



$d=12$ ,  $R^2=1$



# Activity wrap-up (3/3)

- **$R^2$  increases** as the degree increases as there is more and more “freedom” in how the polynomial curve can twist.
- **The fit gets stranger and stranger** (negative values when  $d=12\dots$ )

## Conclusions:

- This is known as **overfitting**: regression curve fits “too closely” to existing datapoints.
- Ends up **not reflecting reality** as too tailored to the dataset we have: poor prediction abilities
- Need to find **new ways of measuring** what a “good fit” is!

How to measure how good a model is at prediction?

# How good is a model at predicting? (1/3)

If we only use the data at hand to evaluate our model, then the model can **overfit to the data**.

**What if I gave you additional datapoints? Could we now evaluate how good the model is at predicting?**

Time	Sales
10	146.4116
35	195.8377
70	241.7297
135	216.1947
163	166.1761
192	104.6491
228	36.80205
302	14.49156

**Yes!**

**Idea:** See how good the model is at predicting sales for these new points by comparing to “real” sales.

⇒ Tells you how good the model is when faced with points it's never seen.

# How good is a model at predicting? (2/3)

We predict the sales for these new time points using our **linear regression model**.

Time	Sales
10	146.4116
35	195.8377
70	241.7297
135	216.1947
163	166.1761
192	104.6491
228	36.80205
302	14.49156

$$\text{Root Mean Squared Error (RMSE)} = \sqrt{\frac{1}{n} \cdot ((pred_1 - actual_1)^2 + \dots + (pred_n - actual_n)^2)}$$

# How good is a model at predicting? (3/3)



We can do this for each model we looked at.

Degree	RMSE
1	51.18
2	54.5
3	18.6
5	21.33
8	20.86
12	93.26

- The best model from this table is degree =3
- This ties in with what our intuition would have had us pick!


Here, I gave you new datapoints for this to work. But how would we proceed if we didn't have new datapoints?

**Idea: When you first get your dataset, split it up and only build your model on part of it! Use the other part to evaluate it.**

# Training, validation, and test sets

# Training, validation, and test sets (1/2)

We in fact divide the existing data into **three**.




- The **training set** serves to **build your model**.
- The **validation set** serves to **select between models**.
- The **test set** is used just once to **give an indication as to how the chosen model will perform**.


# Training, validation, and test sets (2/2)

## Example: Polynomial regression

1. Use the **training set** to come up with polynomial regressors with different degrees.



$d=3$



$d=8$

2. Use the **validation set** to pick which degree is the best, e.g.  $d = 3$ .
3. Use the **testing set** to evaluate how well the model you picked would perform on new data.



# These concepts in Python

We use scikit again.

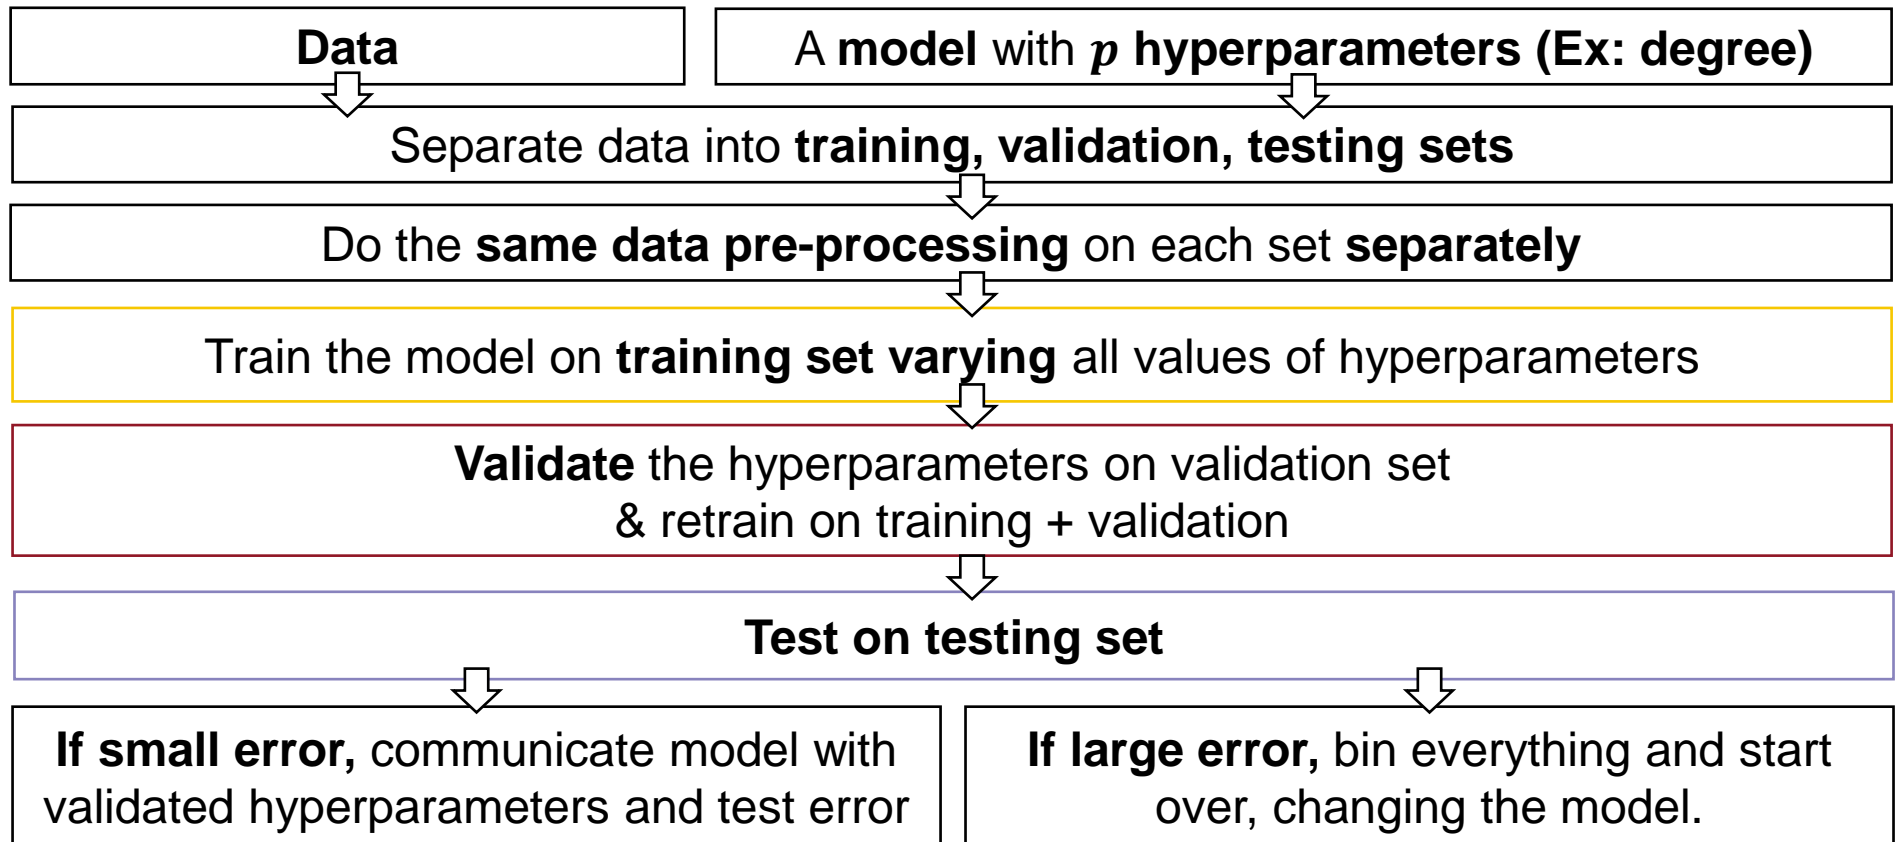
- How to **divide** a dataset into **two**:

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(Xdata, ydata, test_size=0.33)
```

- How to **divide** a dataset into **three**: use the function above twice (see homework)
- How to **compute the RMSE** between real values and predicted values (example here with the linear regression model)

```
y_pred_val=lm.predict(Xval)
from sklearn.metrics import mean_squared_error
mean_squared_error(Yval,y_pred_val)**(1/2)
```

# Supervised learning process



- Ideally: have different validation sets for each parameter so as not to **overfit the validation set**
- **Requires a lot of data /computation power**

# Wrap-up & Next time

Today, we:

- **Reviewed linear and polynomial regression**
- Understood the concept of **overfitting**
- Understood the purpose of **training/testing/validation sets**
- Saw **how to approach** supervised learning problems such as regression

Next time:

- Pitfalls of Machine Learning + Starting on Classification
- Mini-quiz to do + homework



The Business School  
for the World®

EUROPE

|

ASIA

|

MIDDLE EAST

|

NORTH AMERICA